



# Computer are dumb

(And why you care.)

Titus Brown

6/10/11



# Outline

- Algorithms and scaling
- Heuristics in computation
- Some thoughts on hardware
- Are you right? (Or at least not wrong?)
- Whose fault is it that you're probably wrong? (RCR training!)



## Consider: finding SNPs.

Given: reference genome, sequence reads, mapping.

The mapping contains a list of reads, mapped locations within reference, and the location of differences.

How can we find all single-nucleotide variation?



# Approach one: by genome

```
for location in genome:  
    reference = genome[location]  
    bases = get_overlapping(location)  
    for base in bases:  
        if base != reference:  
            # count SNP
```



## Approach two: by read

```
for read in mapped_reads:  
    for differences in read:  
        # count SNP
```

# Approach one: by genome

```
for location in genome:
    reference = genome[location]
    bases = get_overlapping(location)
    for base in bases:
        if base != reference:
            # count SNP
```

How does this algorithm scale?

Imagine:

- increasing size of genome

- increasing number of reads



## Approach two: by read

```
for read in mapped_reads:  
    for differences in read:  
        # count SNP
```

How does *this* algorithm scale?

# Scaling and Big-O notation

- The first approach scales with both the size of the genome and the number of reads:

$$t \sim O(N * M)$$

- The second approach scales with just the number of reads:

$$t \sim O(M)$$



# Scaling and Big-O notation

- The first approach scales with both the size of the genome and the number of reads:

$$t \sim O(N * M)$$

- why would you want this??

- The second approach scales with just the number of reads:

$$t \sim O(M)$$

# What about a different problem?

- I am interested in locations X,Y, and Z.
- Give me all SNPs at or near those locations.

```
for location in list_of_locations:  
    reference = genome[location]  
    bases = get_overlapping(location)  
    for base in bases:  
        if base != reference:  
            # count SNP
```



## Important note

- *Algorithm* scaling is independent of the actual time it takes to run.
- Scaling tells you how time-to-run scales as the problem size changes, nothing more.



## Easy-to-check vs easy-to-find

Given a number, factor it into only prime numbers.

*This is hard.*

Given a set of prime numbers, verify that they multiply to yield a particular number.

*This is easy.*

# Easy-to-check vs easy-to-find, #2

Suppose:

50 dorm rooms, two students per room

100 students can be admitted, of 400 total

Dean has list of students that cannot be paired.

It is *easy to check* any particular list of student/room combinations for validity.

In general, it is *extremely hard* to quickly find a guaranteed solution.



# Heuristics

- “Heuristics” are short cuts that usually work (but occasionally go horribly wrong).
- Not all problems are amenable.
  - Prime numbers? No good, *fast* short cut.
  - Housing? Sure – start with a random solution, eliminate one of each pair that conflicts, until you find a non-conflict..
- Heuristics rely on assumptions about the specific type of problem you’re going to tackle, and don’t always work.
  - If the Dean is evil, he can construct a list of incompatible roommates that breaks your process.
  - Or he can just gives you a really long list of incompatible roommates.

# Example: BLAST

BLASTN filters sequences for exact matches between “words” of length  $W$ :

```
GAGGGTATGACGATATGGCGATGGAC
||x|||||x|||||||x|x||x
GAcGGTATcACGATATGGCGgT-Gag
```

This results in a  $O(n \log n)$  algorithm.

## Example: BLAST

...but what about pathological situations?

```
GAGGGTATGACGATATGGCGATGGAC
||x|||||x|||||x|||||x|x||x
GAcGGTATcACGATGTGGCGgT-Gag
```

This will not be scored as a match, because  
BLAST *only* scores matches with a core  
“seed” match of 11 bases.






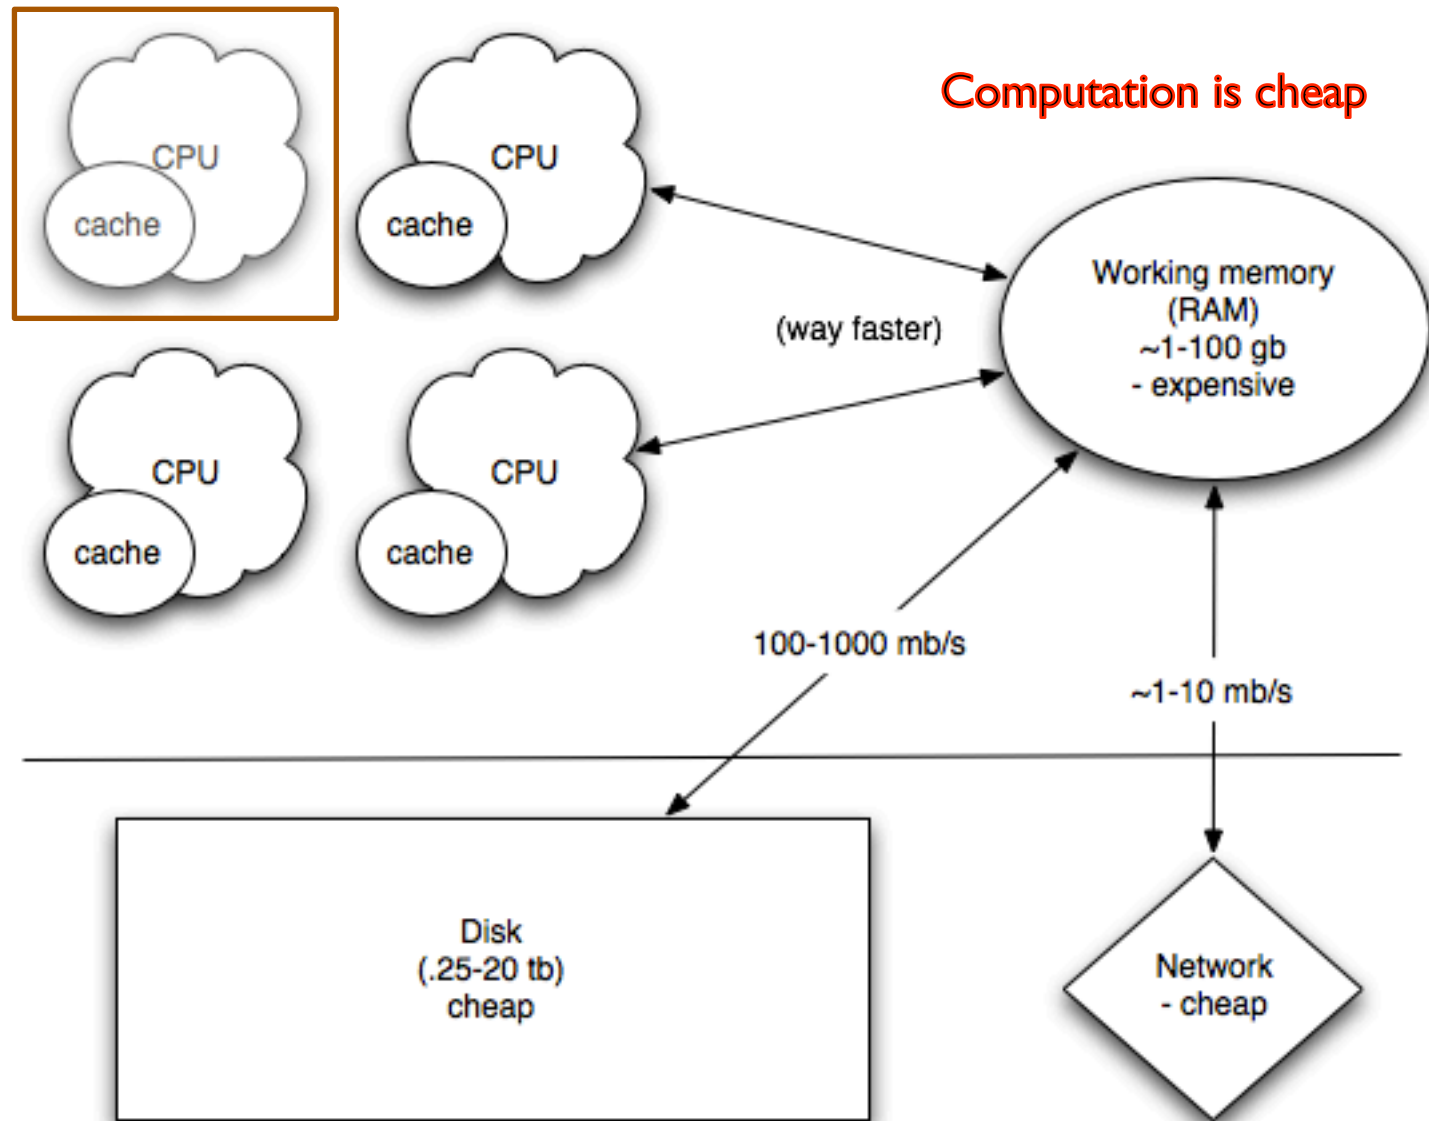
# Heuristics

- Often take advantage of “fast” operations on computers.
- What operations are fast are governed by
  - Algorithmic considerations / fundamental CS
  - Hardware and software implementations

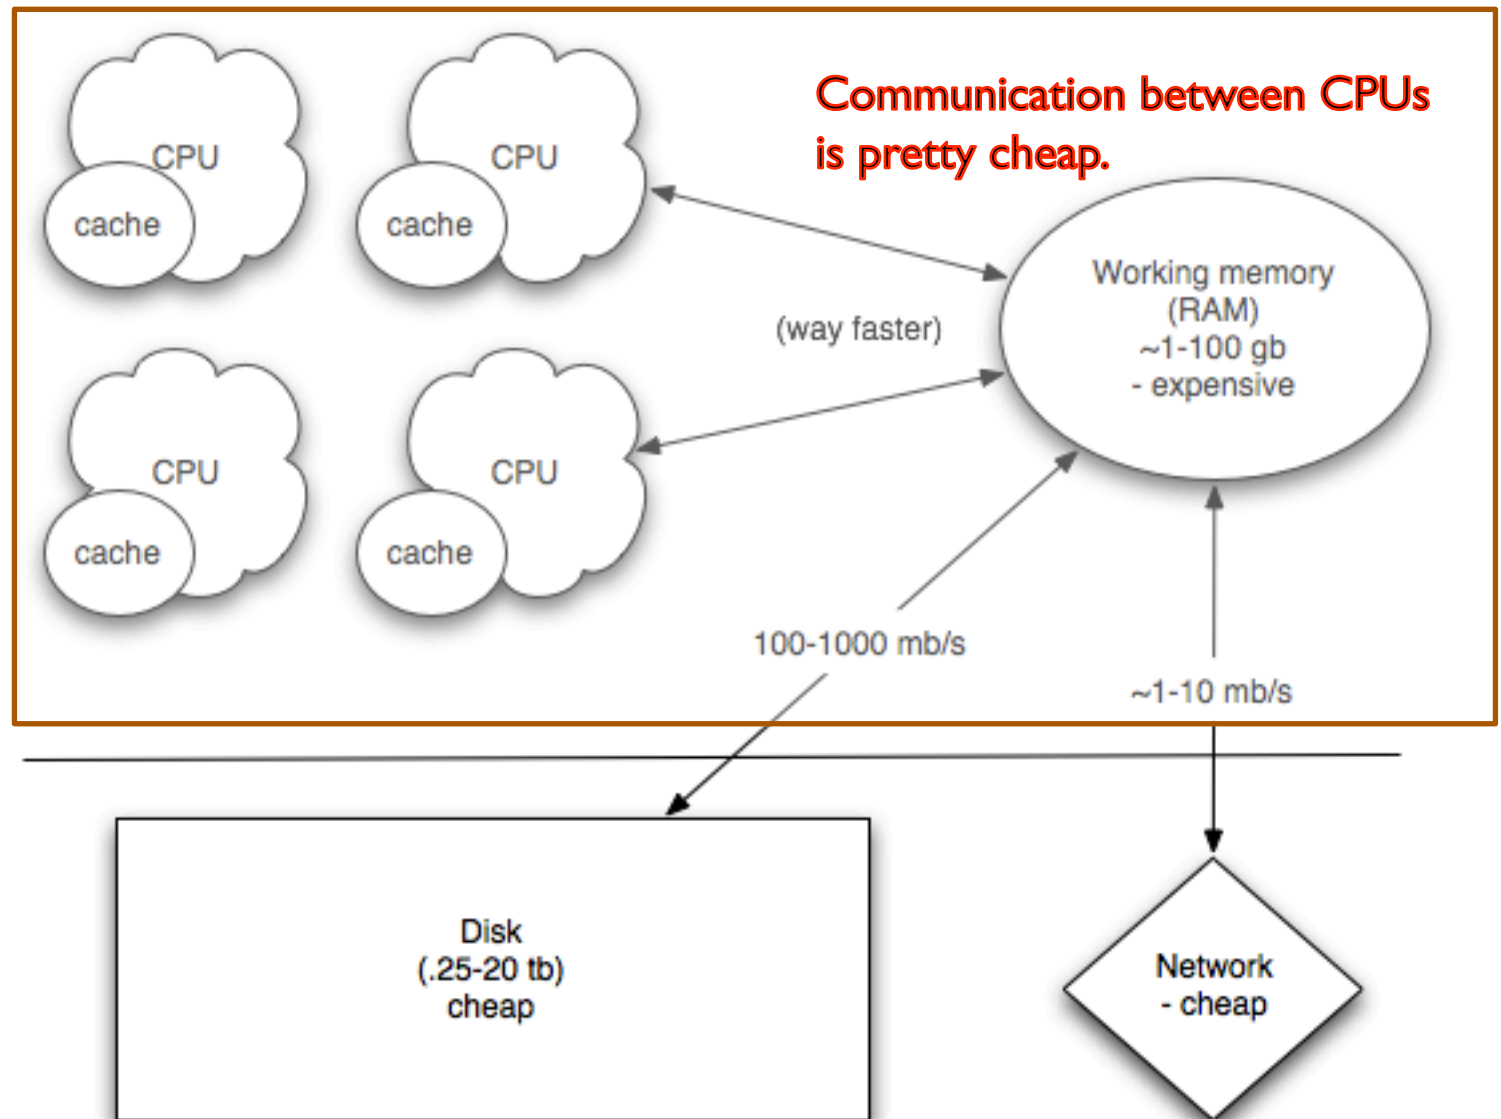
# Computer architecture



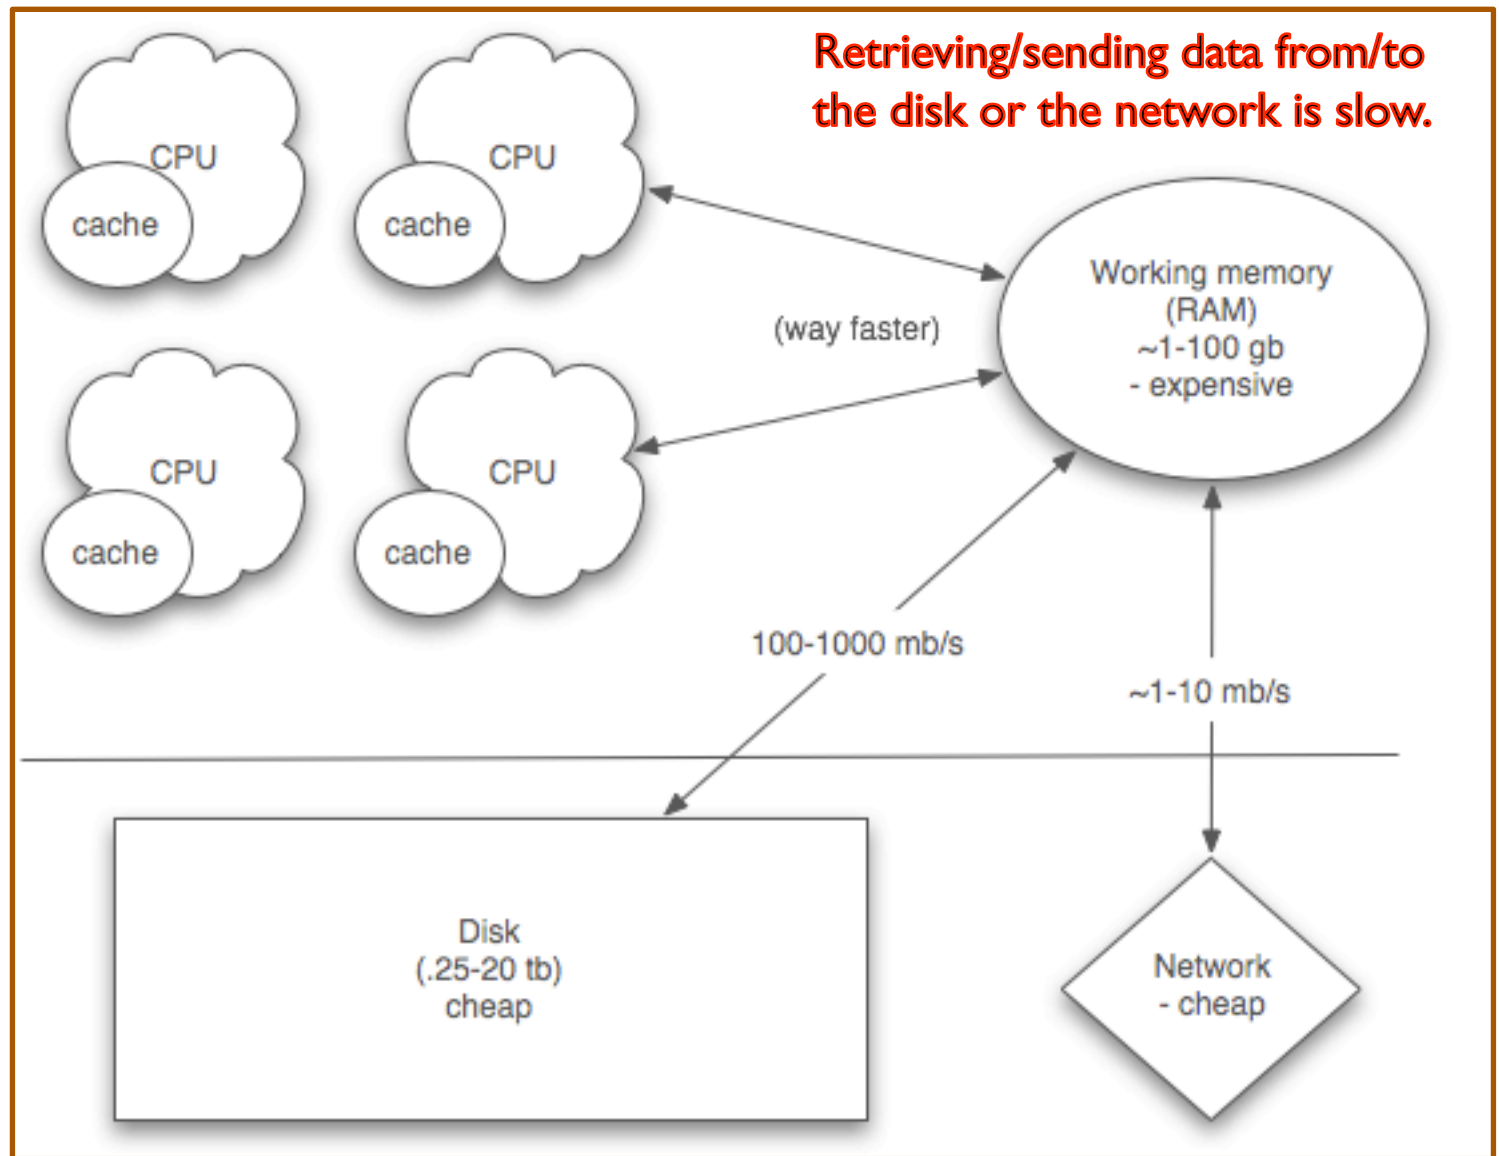
# Computer architecture



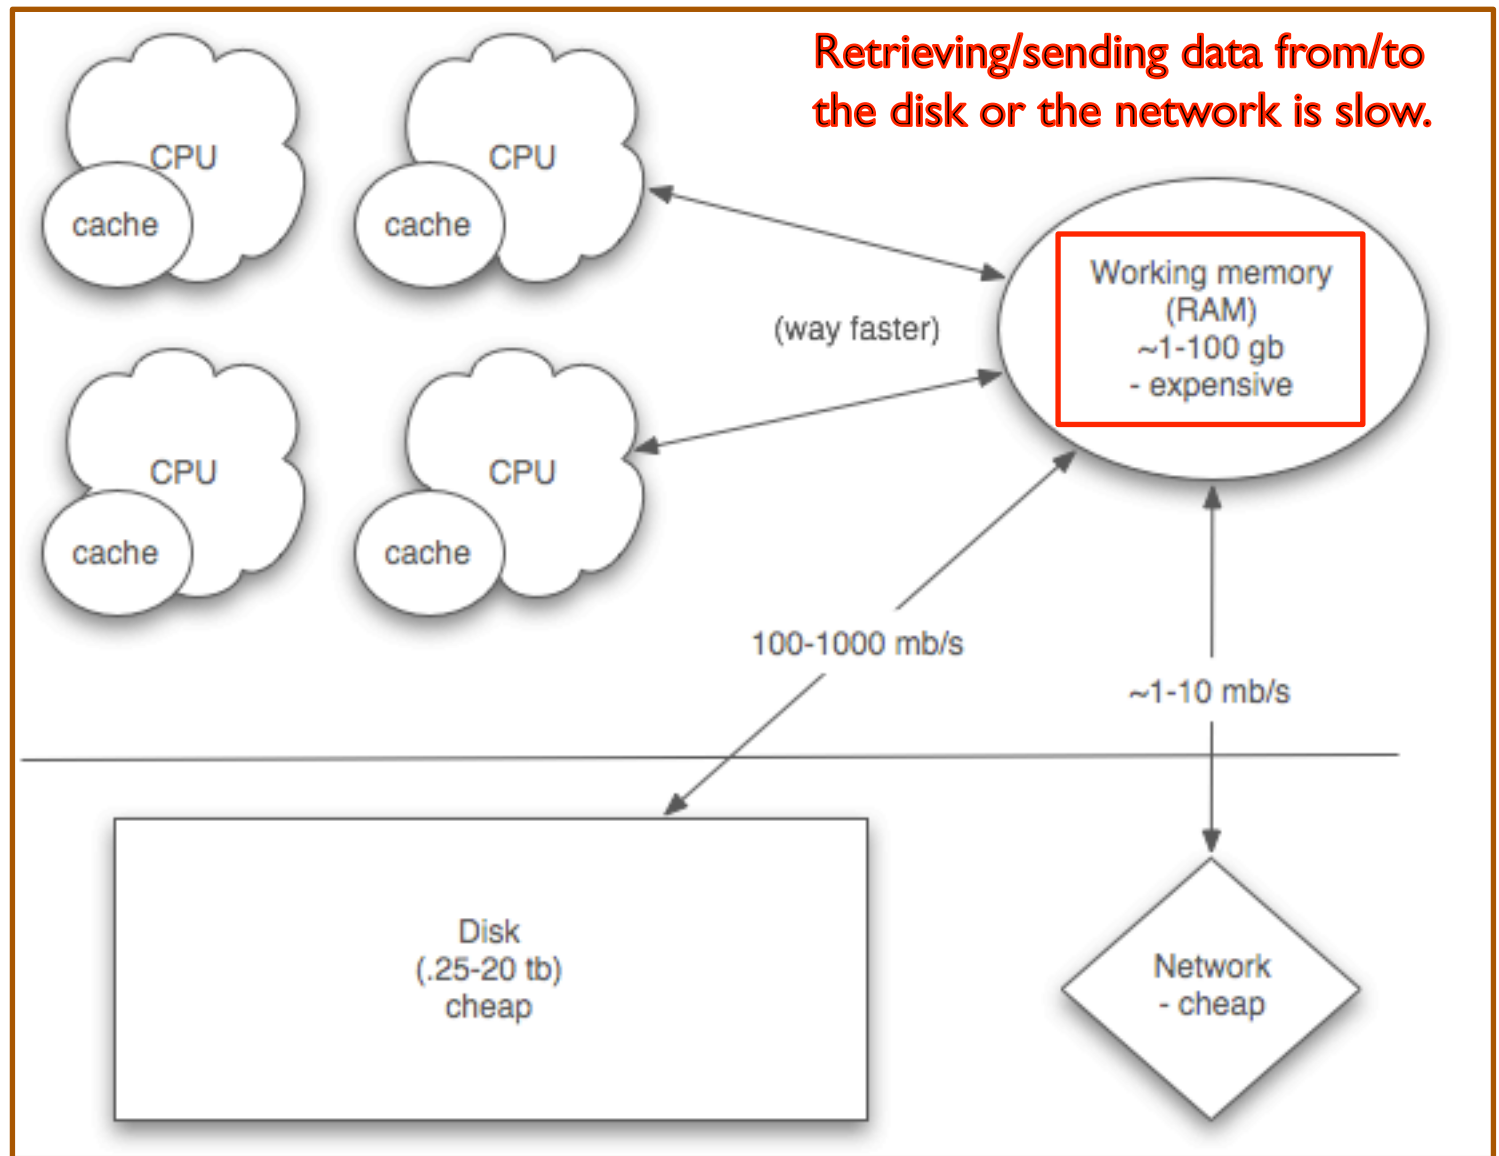
# Computer architecture



# Computer architecture



# Computer architecture

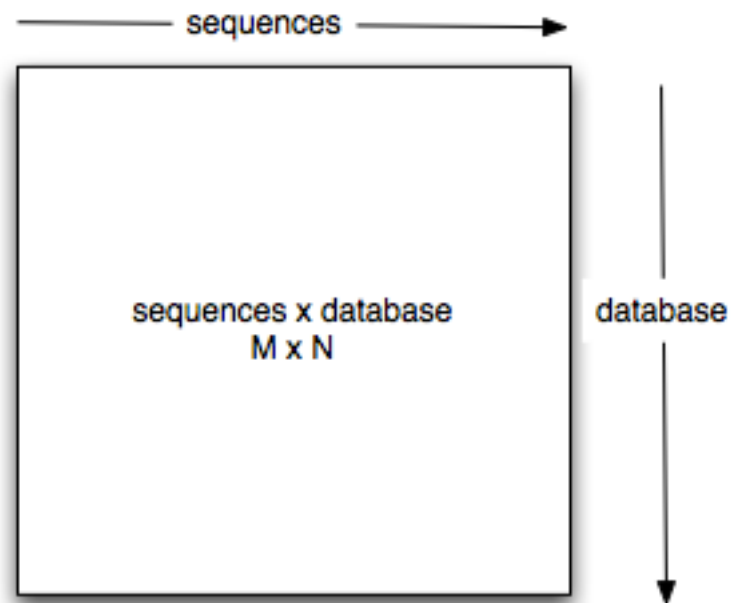




# Questions to ask


- Can I split my problem up into small chunks?  
(because, if so, I can use more than one computer effectively.)
- How does my computation scale?
- How does my memory use scale?

# Sequence comparison: $n^2$





...but “embarrassingly parallel”





Mapping is embarrassingly parallel.

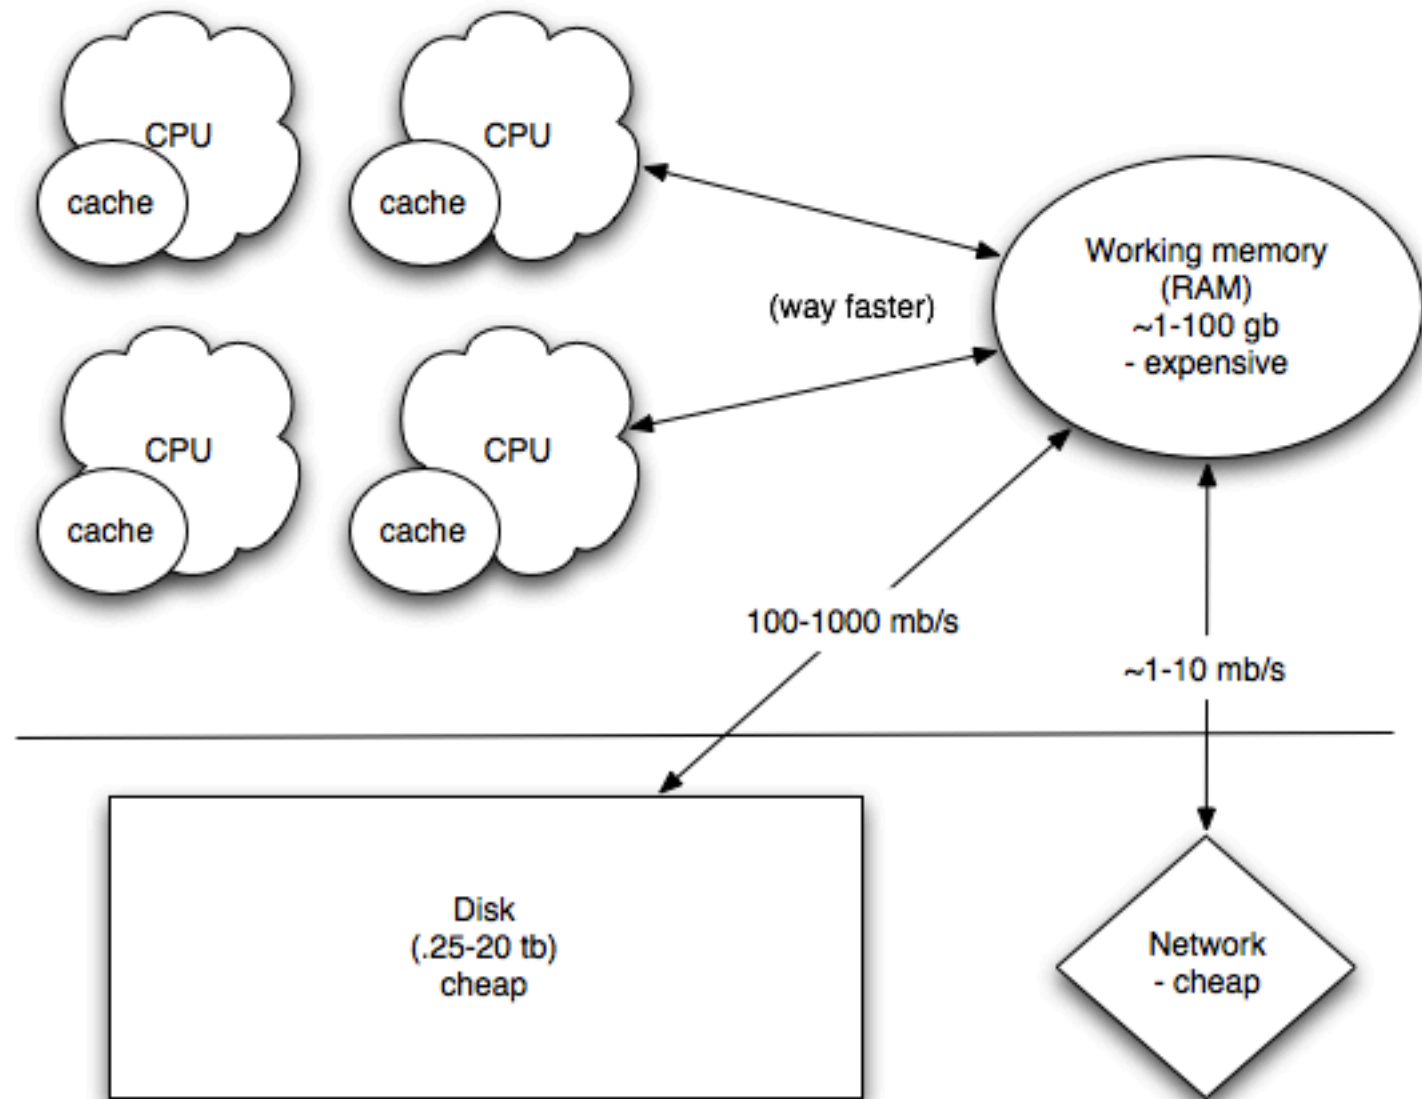
You need to calculate individual overlaps.



# Assembly is not.

You need to calculate *all overlaps*.

Communication between CPUs is slow; this is main factor in splitting up tasks.





None of this is the #1 problem you will face with bioinformatics.

Here is the #1 problem:

How would you know if your answer was right or wrong?



None of this is the #1 problem you will face with bioinformatics.

Here is the #1 problem:

How would you know if your answer was right or wrong?

**If you can't answer this question, then what's the point?**



# Controls

- Just as with experiments, you can put negative and positive controls in your bioinformatics.
- e.g. with BLAST,
  - Do you see expected matches with the parameters and database you're using?
- Positive controls are often easier than negative, in “discovery-driven” science...



# Internal controls

- Molecules and sequences for which you have expectations.
- “I know this gene comes up, based on qPCR. I expect to see it in my mRNAseq.”





# External controls

- Does the whole process work?
- “I can reproduce what this other person/ lab did, with their data, when I use my own software.”
- This is much more rarely done...

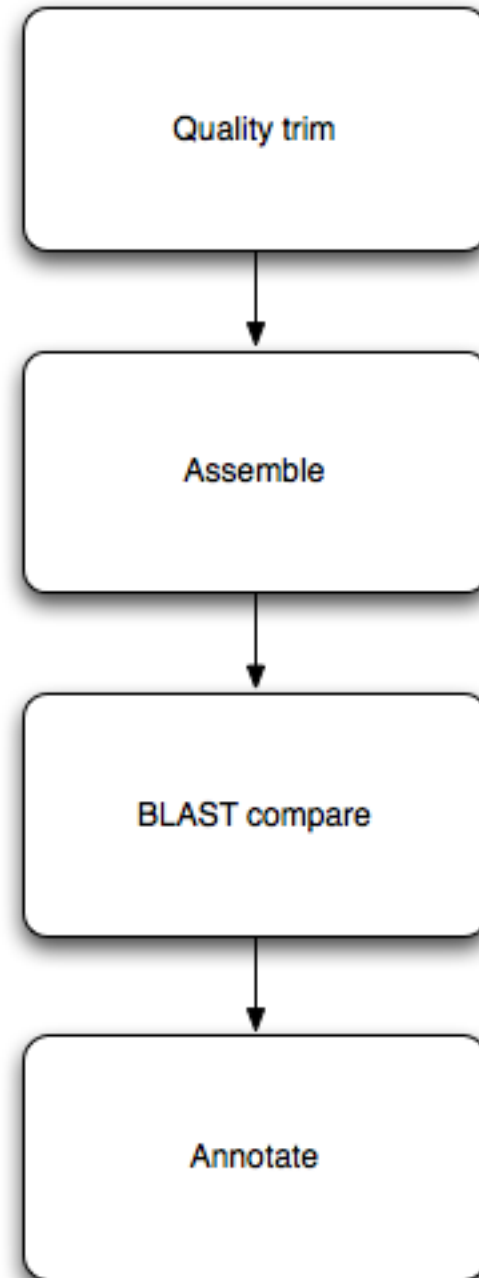


# Black box nature of algorithms

- When you listen to a computational biologist explain their clever algorithm...
- ...it's a mistake to think that they necessarily know what's going on.
- Software is full of bugs and unintended consequences.

# Pipelines

- Each step can be understood, and tested/controlled individually.
- Each step is re-usable! Just need to figure out input/output formats.
- Automate, automate, automate.





# Tracking the process

- How do you know that your software *today* is doing the same thing it was doing last month? Or last year? Or in the hands of that other graduate student?
- There are tools and processes for dealing with this
  - Version control (think “change tracking” in Word)
  - Automated testing (think “automated positive/negative controls”)



# Replication vs Reproducibility

- Replication: identical results.
  - Same parameters  $\Rightarrow$  same results
- Reproducibility: similar results.
  - Similar parameters  $\Rightarrow$  similar results
  - (What's “similar”?)
- Corroboration:
  - Similar results seen in a different system



# What should our standards be?

- For *scientists*, reproducibility is extremely important.
- Replication ... less so. It's very challenging to *exactly* replicate a given experimental situation.
- But, there is a pragmatic reason to think about replication, too.



# Replication in computational science

- We have spent mucho time making sure that computers do the same thing *every time*, at the micro level.
- If you observe *unplanned* variation in a computational system, then:
  - You either are using one of the approximate subsystems, like floating point;
  - **Or you have a bug.**



# Replication in computational science

- So, *proximate to an analysis*, you should be able to exactly replicate a particular result from a computational analysis.
- Then, changing
  - Data sets
  - Hardware
  - Underlying software
- ...may result in observed differences.





# Concluding thoughts

- Every step of the process needs to be critically thought about and controlled.
- Choice of algorithms can be important, but depends on your problem: the convenient tool in your toolbox may not be well suited to your problem.